



Assignment 2

Part 1: Fish Market Regression

Q1. Linear Regression (15 points)

Import Libraries and Load Dataset

```
In [5]: import pandas as pd
import numpy as np

from sklearn.model_selection import train_test_split
import statsmodels.api as sm

import matplotlib.pyplot as plt

df = pd.read_excel("Fish.xls")
print(df.head())
```

	Species	Weight	Length1	Length2	Length3	Height	Width
0	Bream	242.0	23.2	25.4	30.0	11.5200	4.0200
1	Bream	290.0	24.0	26.3	31.2	12.4800	4.3056
2	Bream	340.0	23.9	26.5	31.1	12.3778	4.6961
3	Bream	363.0	26.3	29.0	33.5	12.7300	4.4555
4	Bream	430.0	26.5	29.0	34.0	12.4440	5.1340

```
In [6]: print(df.columns)

Index(['Species', 'Weight', 'Length1', 'Length2', 'Length3', 'Height',
      'Width'],
      dtype='object')
```

Define Features and Target

```
In [8]: X = df[['Length1', 'Length2', 'Length3', 'Height', 'Width']]
y = df['Weight']

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.30, random_state=42
)

print("Train size:", X_train.shape, " Test size:", X_test.shape)
```

Train size: (111, 5) Test size: (48, 5)

Fit Linear Regression (OLS)

```
In [24]: X_train_sm = sm.add_constant(X_train) # Add intercept
```

```
ols = sm.OLS(y_train, X_train_sm).fit()

print(ols.summary())
```

OLS Regression Results

Dep. Variable:	Weight	R-squared:	0.888			
Model:	OLS	Adj. R-squared:	0.883			
Method:	Least Squares	F-statistic:	167.0			
Date:	Wed, 18 Feb 2026	Prob (F-statistic):	2.66e-48			
Time:	14:54:10	Log-Likelihood:	-689.14			
No. Observations:	111	AIC:	1390.			
Df Residuals:	105	BIC:	1407.			
Df Model:	5					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]

const	-530.2590	36.848	-14.391	0.000	-603.321	-457.197
Length1	32.6915	52.628	0.621	0.536	-71.660	137.043
Length2	25.0097	52.608	0.475	0.635	-79.302	129.321
Length3	-29.7910	19.315	-1.542	0.126	-68.089	8.507
Height	23.8780	10.262	2.327	0.022	3.530	44.226
Width	15.9742	23.600	0.677	0.500	-30.821	62.769
=====						
Omnibus:	8.469	Durbin-Watson:	1.983			
Prob(Omnibus):	0.014	Jarque-Bera (JB):	8.353			
Skew:	0.562	Prob(JB):	0.0154			
Kurtosis:	3.737	Cond. No.	342.			
=====						

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Estimated coefficients, R-squared term, and the 95% confidence interval

```
In [25]: print("Coefficients (Least Squares Estimates):")
print(ols.params)
```

```
Coefficients (Least Squares Estimates):
const      -530.258960
Length1     32.691486
Length2     25.009744
Length3    -29.791039
Height      23.878024
Width       15.974163
dtype: float64
```

```
In [103]: print("\nStandard Errors:")
print(ols.bse)
```

```
Standard Errors:
const      36.847650
Length1    52.628116
Length2    52.607849
Length3    19.314951
Height     10.262141
Width      23.600338
dtype: float64
```

```
In [27]: print("\nR-squared:", ols.rsquared)
        print("Adjusted R-squared:", ols.rsquared_adj)
```

```
R-squared: 0.8882734450295612
Adjusted R-squared: 0.8829531328881117
```

```
In [28]: print("\n95% Confidence Intervals:")
        print(ols.conf_int(alpha=0.05))
```

```
95% Confidence Intervals:
              0              1
const  -603.321043 -457.196877
Length1 -71.660346 137.043319
Length2 -79.301903 129.321391
Length3 -68.089019  8.506941
Height   3.530094  44.225954
Width   -30.820946  62.769273
```

Finding dependence between the length and weight

```
In [29]: print("p-values:")
        print(ols.pvalues)
```

```
p-values:
const      1.380947e-26
Length1    5.358280e-01
Length2    6.354898e-01
Length3    1.259892e-01
Height     2.189513e-02
Width      4.999822e-01
dtype: float64
```

```
In [20]: corr = df[['Weight', 'Length1', 'Length2', 'Length3']].corr()
        print(corr['Weight'])
```

```
Weight      1.000000
Length1     0.915712
Length2     0.918618
Length3     0.923044
Name: Weight, dtype: float64
```

Is there any dependence between the length and weight of the fish?

Yes, there is strong dependence between fish length and weight.

The correlation coefficients between Weight and the length variables (Length1 = 0.9157, Length2 = 0.9186, Length3 = 0.9230) are very high, indicating a strong positive linear relationship. This means that as fish length increases, fish weight also increases significantly.

Additionally, the regression model has a high R^2 value of 0.888, showing that about 88.8% of the variation in weight is explained by the size-related variables. Although the individual p-values for the length variables are not significant due to multicollinearity, the overall analysis clearly indicates a strong dependence between fish length and weight.

Q2. Ridge and Lasso Regression

Fit Ridge and Lasso Model and Report Coefficients

```
In [22]: from sklearn.linear_model import Ridge, Lasso

# Ridge Regression
ridge = Ridge(alpha=1.0)
ridge.fit(X_train, y_train)

ridge_coef = pd.Series(ridge.coef_, index=X.columns)
print("Ridge Coefficients:")
print(ridge_coef)

# Lasso Regression
lasso = Lasso(alpha=1.0, max_iter=10000)
lasso.fit(X_train, y_train)

lasso_coef = pd.Series(lasso.coef_, index=X.columns)
print("\nLasso Coefficients:")
print(lasso_coef)
```

```
Ridge Coefficients:
Length1    30.970924
Length2    25.243218
Length3   -28.474155
Height     23.229551
Width      16.819504
dtype: float64
```

```
Lasso Coefficients:
Length1    36.742802
Length2    16.515789
Length3   -25.291690
Height     22.379919
Width      17.703329
dtype: float64
```

Least Impactful Attribute Based on Ridge and Lasso Regression

Based on the Ridge and Lasso results, Width is the least impactful attribute in predicting fish weight. In the Ridge model, Width has the smallest absolute coefficient (16.82), and in the Lasso model it also has one of the smallest coefficients (17.70). Since smaller coefficient magnitudes indicate weaker influence on the target variable, Width contributes the least to predicting fish weight.

Q3. Model Testing

```
In [31]: from sklearn.metrics import mean_squared_error, r2_score

# OLS
X_test_sm = sm.add_constant(X_test)
y_pred_ols = ols.predict(X_test_sm)

mse_ols = mean_squared_error(y_test, y_pred_ols)
r2_ols = r2_score(y_test, y_pred_ols)

# Ridge
y_pred_ridge = ridge.predict(X_test)

mse_ridge = mean_squared_error(y_test, y_pred_ridge)
r2_ridge = r2_score(y_test, y_pred_ridge)

# Lasso
y_pred_lasso = lasso.predict(X_test)

mse_lasso = mean_squared_error(y_test, y_pred_lasso)
r2_lasso = r2_score(y_test, y_pred_lasso)

# Summary
results = pd.DataFrame({
    "Model": ["OLS", "Ridge", "Lasso"],
    "MSE": [mse_ols, mse_ridge, mse_lasso],
    "R2": [r2_ols, r2_ridge, r2_lasso]
})
print(results.sort_values(by="MSE"))
print("\nBest by MSE:", results.loc[results["MSE"].idxmin(), "Model"])
print("Best by R2:", results.loc[results["R2"].idxmax(), "Model"])
```

	Model	MSE	R2
2	Lasso	16203.776733	0.867718
0	OLS	16217.627414	0.867605
1	Ridge	16239.351882	0.867427

Best by MSE: Lasso

Best by R2: Lasso

Analysis

Based on the test dataset results, Lasso regression performs the best, achieving the lowest MSE (16203.78) and the highest R^2 (0.8677). Although the differences among the three models are small, Lasso slightly outperforms OLS and Ridge, likely due to its stronger regularization effect, which helps reduce variance and improve generalization. Overall, all three models demonstrate similar performance, indicating a strong linear relationship between fish dimensions and weight.

Part 2: Social Network Logistic Regression

(a) Age and Salary Model

Load Dataset

```
In [59]: df = pd.read_csv("Social_Network_ads.csv")

print(df.head())
print(df.info())
```

```
   User ID  Gender  Age  EstimatedSalary  Purchased
0  15624510   Male   19             19000           0
1  15810944   Male   35             20000           0
2  15668575  Female   26             43000           0
3  15603246  Female   27             57000           0
4  15804002   Male   19             76000           0
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   User ID                400 non-null   int64   
1   Gender                 400 non-null   object  
2   Age                    400 non-null   int64   
3   EstimatedSalary        400 non-null   int64   
4   Purchased              400 non-null   int64   
dtypes: int64(4), object(1)
memory usage: 15.8+ KB
None
```

Select Features

```
In [63]: X = df[['Age', 'EstimatedSalary']]
         y = df['Purchased']
```

Train-Test Split

```
In [64]: X_train, X_test, y_train, y_test = train_test_split(
        X, y, test_size=0.30, random_state=42
    )

    print("X_train shape:", X_train.shape)
    print("X_test shape:", X_test.shape)
    print("y_train shape:", y_train.shape)
    print("y_test shape:", y_test.shape)
```

```
X_train shape: (280, 2)
X_test shape: (120, 2)
y_train shape: (280,)
y_test shape: (120,)
```

Feature Scaling

```
In [65]: from sklearn.preprocessing import StandardScaler
        from sklearn.linear_model import LogisticRegression

        scaler = StandardScaler()

        X_train = scaler.fit_transform(X_train)
        X_test = scaler.transform(X_test)
```

```
In [52]: print(X_train[:5])

[[-0.72378527 -0.69308934 -0.75425342 -0.63284698 -0.58010567]
 [-0.83421159 -0.81507103 -0.91108538 -0.95932281 -1.04308495]
 [ 0.75191187  0.73315817  0.65723425  0.49477347  1.44315024]
 [ 0.2198578  0.21708177  0.41327342  1.36998948  0.19796686]
 [ 0.40055541  0.3859795  0.22158991  0.28138565  0.92913387]]
```

Train Logistic Regression Model

```
In [66]: log_reg = LogisticRegression()
        log_reg.fit(X_train, y_train)
```

Out[66]:

LogisticRegression i ?		
Parameters		
	penalty	'l2'
	dual	False
	tol	0.0001
	C	1.0
	fit_intercept	True
	intercept_scaling	1
	class_weight	None
	random_state	None
	solver	'lbfgs'
	max_iter	100
	multi_class	'deprecated'
	verbose	0
	warm_start	False
	n_jobs	None
	l1_ratio	None

Make Predictions

```
In [69]: y_pred = log_reg.predict(X_test)
y_prob = log_reg.predict_proba(X_test)[:, 1]
```

Performance Metrics

```
In [71]: from sklearn.metrics import (
accuracy_score, precision_score, recall_score,
f1_score, roc_curve, roc_auc_score
)
```

```
In [73]: accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

print("Accuracy:", accuracy)
print("Precision:", precision)
```




```
print("Recall:", recall)
print("F1 Score:", f1)
```

Accuracy: 0.85
Precision: 0.9393939393939394
Recall: 0.6595744680851063
F1 Score: 0.775

```
In [76]: fpr, tpr, thresholds = roc_curve(y_test, y_prob)
roc_auc = roc_auc_score(y_test, y_prob)

plt.figure(figsize=(6,5))
plt.plot(fpr, tpr, label="ROC Curve (AUC = {:.2f})".format(roc_auc))
plt.plot([0,1], [0,1], linestyle='--')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve")
plt.legend()
plt.show()
```



(b) Gender, Age, and Salary Model

Update the features

```
In [79]: # Encode Gender: Female=0, Male=1
df["Gender_bin"] = df["Gender"].map({"Female": 0, "Male": 1})

X = df[["Gender_bin", "Age", "EstimatedSalary"]]
y = df["Purchased"]
```

Split the data (70/30)

```
In [97]: X_train, X_test, y_train, y_test = train_test_split(
        X, y, test_size=0.30, random_state=42, stratify=y
    )

# Keep gender labels for test rows so we can evaluate male/female separately
gender_test = df.loc[X_test.index, "Gender"]
print(gender_test)
```

```
300    Female
195     Male
215    Female
276     Male
395    Female
...
342    Female
66     Male
109    Female
335    Female
224    Female
Name: Gender, Length: 120, dtype: object
```

Scale Age and Salary

```
In [82]: scaler = StandardScaler()
X_train_scaled = X_train.copy()
X_test_scaled = X_test.copy()

X_train_scaled[["Age", "EstimatedSalary"]] = scaler.fit_transform(X_train[["Age", "EstimatedSalary"]])
X_test_scaled[["Age", "EstimatedSalary"]] = scaler.transform(X_test[["Age", "EstimatedSalary"]])
```

Train model

```
In [83]: model = LogisticRegression()
model.fit(X_train_scaled, y_train)
```

Out[83]:

LogisticRegression i ?		
Parameters		
	penalty	'l2'
	dual	False
	tol	0.0001
	C	1.0
	fit_intercept	True
	intercept_scaling	1
	class_weight	None
	random_state	None
	solver	'lbfgs'
	max_iter	100
	multi_class	'deprecated'
	verbose	0
	warm_start	False
	n_jobs	None
	l1_ratio	None

Make Predictions

```
In [93]: y_pred = model.predict(X_test_scaled)
y_prob = model.predict_proba(X_test_scaled)[: , 1]
# Put predictions into Series aligned with X_test index
y_pred_s = pd.Series(y_pred, index=X_test.index)
y_prob_s = pd.Series(y_prob, index=X_test.index)

print(y_pred)
print(y_pred_s)
```

```

[1 0 1 0 0 0 0 0 0 1 0 0 0 1 0 0 1 1 0 0 0 0 0 0 0 0 1 0 0 0 1 0 1 0 1 0 0
 0 0 0 0 1 0 1 1 0 0 1 1 0 1 1 1 1 0 1 1 0 1 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 1
 1 0 0 1 0 0 0 0 0 1 0 0 0 0 0 0 1 1 1 0 0 1 0 0 0 0 0 0 0 0 1 0 1 0 0 0 0 0
 1 0 0 0 0 0 0 0 0]
300    1
195    0
215    1
276    0
395    0
...
342    0
66     0
109    0
335    0
224    0
Length: 120, dtype: int64

```

Function for Computing Metrics

```

In [94]: def compute_group_metrics(group_name, mask):
          yt = y_test[mask]
          yp = y_pred_s[mask]
          ypr = y_prob_s[mask]

          metrics = {
              "Group": group_name,
              "n": int(mask.sum()),
              "Accuracy": accuracy_score(yt, yp),
              "Precision": precision_score(yt, yp, zero_division=0),
              "Recall": recall_score(yt, yp, zero_division=0),
              "F1": f1_score(yt, yp, zero_division=0),
              "AUC": roc_auc_score(yt, ypr) if len(np.unique(yt)) > 1 else np.nan
          }
          return metrics, yt, ypr

```

```

In [99]: male_mask = (gender_test == "Male")
          female_mask = (gender_test == "Female")
          print(male_mask)

```

```

300    False
195     True
215    False
276     True
395    False
...
342    False
66     True
109    False
335    False
224    False
Name: Gender, Length: 120, dtype: bool

```

Compute Merics

```
In [100... male_metrics, y_male, p_male = compute_group_metrics("Male", male_mask)
female_metrics, y_female, p_female = compute_group_metrics("Female", female_ma
results = pd.DataFrame([male_metrics, female_metrics])
print(results)
```

	Group	n	Accuracy	Precision	Recall	F1	AUC
0	Male	65	0.815385	0.714286	0.555556	0.625	0.900709
1	Female	55	0.836364	0.900000	0.720000	0.800	0.926667

ROC Curve


```
In [101... plt.figure(figsize=(7,6))

# Male ROC
if len(np.unique(y_male)) > 1:
    fpr_m, tpr_m, _ = roc_curve(y_male, p_male)
    auc_m = roc_auc_score(y_male, p_male)
    plt.plot(fpr_m, tpr_m, label=f"Male ROC (AUC={auc_m:.2f})")

# Female ROC
if len(np.unique(y_female)) > 1:
    fpr_f, tpr_f, _ = roc_curve(y_female, p_female)
    auc_f = roc_auc_score(y_female, p_female)
    plt.plot(fpr_f, tpr_f, label=f"Female ROC (AUC={auc_f:.2f})")

# Random baseline
plt.plot([0,1], [0,1], linestyle="--", label="Random (AUC=0.50)")

plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curves by Gender (Logistic Regression)")
plt.legend()
plt.show()
```



Findings of performance differential across gender

The results indicate a noticeable performance difference across gender, with the model performing better for females than males. The most significant differences are observed in precision (0.90 for females vs 0.71 for males) and recall (0.72 for females vs 0.56 for males), leading to a substantially higher F1-score for females (0.80 vs 0.63). This suggests that the model is more accurate and reliable in identifying female buyers compared to male buyers. Additionally, the higher AUC for females (0.93 vs 0.90) further confirms stronger classification performance for the female subgroup.